

ClauseGuard

Find the clause. Verify the words.

We help English-reading students and app users locate important ToS clauses and check the original passage.

LOCAL NLP

8 CATEGORIES

SOURCE-LINKED

A PASSAGE TO REVIEW

“We may remove any photos or comments you upload ... without notice.”

Content removal

TEAM • EVERYONE WRITES CODE

Ankan Roy

Operations

Ameer Shaik

Product

Hemanth Reddy

Engineering

Jayakrishna Puttur

Data & evaluation

Niveda Jawahar

Users & research

The hard part is finding the passage that matters.

Target user: an English-reading student or app user reviewing an unfamiliar Terms of Service.

01

A concrete question

Can the service delete my photos without warning?

02

Today's workflow

Read the document or use keyword search, then interpret the clause.

03

Our value hypothesis

Find a relevant passage and category, then check the exact words.

Focused scope

Eight clause categories, including content removal, arbitration and liability. Privacy and data-sharing detection are outside the current model.

A working app, with the evidence beside each flag.

Desktop app + command-line analysis + self-contained HTML export.

ClauseGuard

Find important clauses. Check the original words.

Flags identify categories to review. They do not determine legality or safety. Unflagged text can still matter. Scores are uncalibrated model outputs, not legal-risk probabilities.

Open text file

Load fictional example

Analyze terms

Export analysis

Category: All sentences

Paste English terms

Fictional service terms for a ClauseGuard demonstration

By accessing or using this service, you agree to these terms.

We may change these terms at any time without notice.

We may suspend or terminate your account at our sole discretion.

We may remove any content you submit without giving a reason.

To the maximum extent permitted by applicable law, we are not liable for indirect or consequential damages.

These terms are governed by the laws of the State of California.

You agree that any court proceedings must be brought exclusively in the courts of San Francisco.

All disputes will be resolved through binding arbitration.

You can download a copy of your account information from the settings page.

Our support team is available Monday through Friday.

Select a result to inspect its source

1. No category flagged
2. Contract by using
- 3. Unilateral termination, Unilateral change**
4. Unilateral termination, Content removal
5. Content removal
6. Limitation of liability
7. Choice of law
8. Jurisdiction
9. Arbitration
10. No category flagged
11. No category flagged

We may change these terms at any time without notice.

Unilateral termination

Lets the provider end or suspend the agreement or account.

Model score: 0.841

Unilateral change

Lets the provider change the agreement or service conditions.

Model score: 1.000

1 Load + analyze

Paste English text or open a text file.

2 Select a result


See the original sentence highlighted.

3 Inspect the source

Read the passage before relying on a tag.

A small, reproducible NLP pipeline that runs locally.

Multi-label classification: one sentence can receive more than one category.



UNFAIR-ToS / LexGLUE • OFFICIAL COMPANY-SEPARATED SPLITS




Shipped: word n-grams 1–2 + character n-grams 3–5; fixed threshold > 0.5.

No runtime API calls. The app and evaluation use the same model artifact.

The shipped hybrid model beats our unigram baseline.

Eight-category macro-F1 on the same 2,275 validation sentences; higher is better.



Macro-F1 weights all eight categories equally. These are model comparisons, not weekly gains.

The main failure is too many tags on a sentence.

A category prediction is a review cue; it can be wrong.

FICTIONAL STUDY DOCUMENT C • TASK C1

“We may remove any photos or comments you upload at our discretion and without notice.”

Content removal ✓

Unilateral termination ✕

Unilateral change ✕

49%

of predicted tags lack
the corresponding
annotation
195 of 398 category tags
on validation sentences

Weakest category: arbitration $F1 = 0.323$ • only 9 positive validation examples.

Four users succeeded—but usefulness is still unproven.

North-star metric: the right sentence AND meaning, within 180 seconds.

4

outside
participants

1 in Session 05

3 in Session 06

8 trials total

RECORDED OUTCOME

CLAUSEGUARD

MANUAL

Correct within
180 seconds

4 / 4

4 / 4

Observed median time

51.7 s

70.9 s

Over-tagged C1 task with the app: 2 of 2 correct.

Short fictional documents; tasks differ in difficulty; convenience sample. No causal speed claim.

We fixed measurement and display errors first.

Progress is the running product, the evidence, and the changes made from failures.

Pilot measurement failed

Typing inflated the clock; the user answered a different task.

Changed: stop at the spoken answer; have the user restate the prompt.

Old details stayed visible

Loading a new document left the previous category detail on screen.

Fixed: edits or a new document clear old results until analysis runs.

Source text did the work

All 3 new users read a source result; none used the category filter.

Next: emphasize the source and test how readers interpret extra tags.

TRUST BOUNDARIES

Local processing • no default input logging • no flags ≠ safe

English only • European annotation context • licensed data with attribution

Per-category thresholds look promising in development.

Leave-one-company-out: tune on 9 validation companies, score the 10th, repeat.

EXPERIMENT — NOT SHIPPED

Macro-F1 ↑

0.6333 → 0.6768

Fixed 0.5 • shipped

Held-out-company folds

Tags per 100 sentences ↓

17.5 → 11.0

Fixed 0.5 • shipped

Held-out-company folds

Next release: ship threshold parameters, rerun app tests, then repeat outside-user tasks.

Arbitration remains unstable; final test performance is still unknown.

Persevere. Earn stronger evidence before claiming value.

The pre-registered midpoint rule passes; it is a decision heuristic, not proof of usefulness.

OCT 20 • S07

Ship the fix

Per-category thresholds;
heading-filter trial.

OCT 27–NOV 3

Test understanding

More outside users;
“no flags” checks.

NOV 10–24

Challenge + freeze

Fresh-document checks;
freeze + final test.

DEC 1 • 5 PM ET

Submit + demo

Product, demo, report;
slides in repository.

What could change the plan: worse task success or misleading tags → change the interface or job;
a failed model gate or >25-point success deficit → consider a pivot.

Find the clause. Verify the words.

Questions